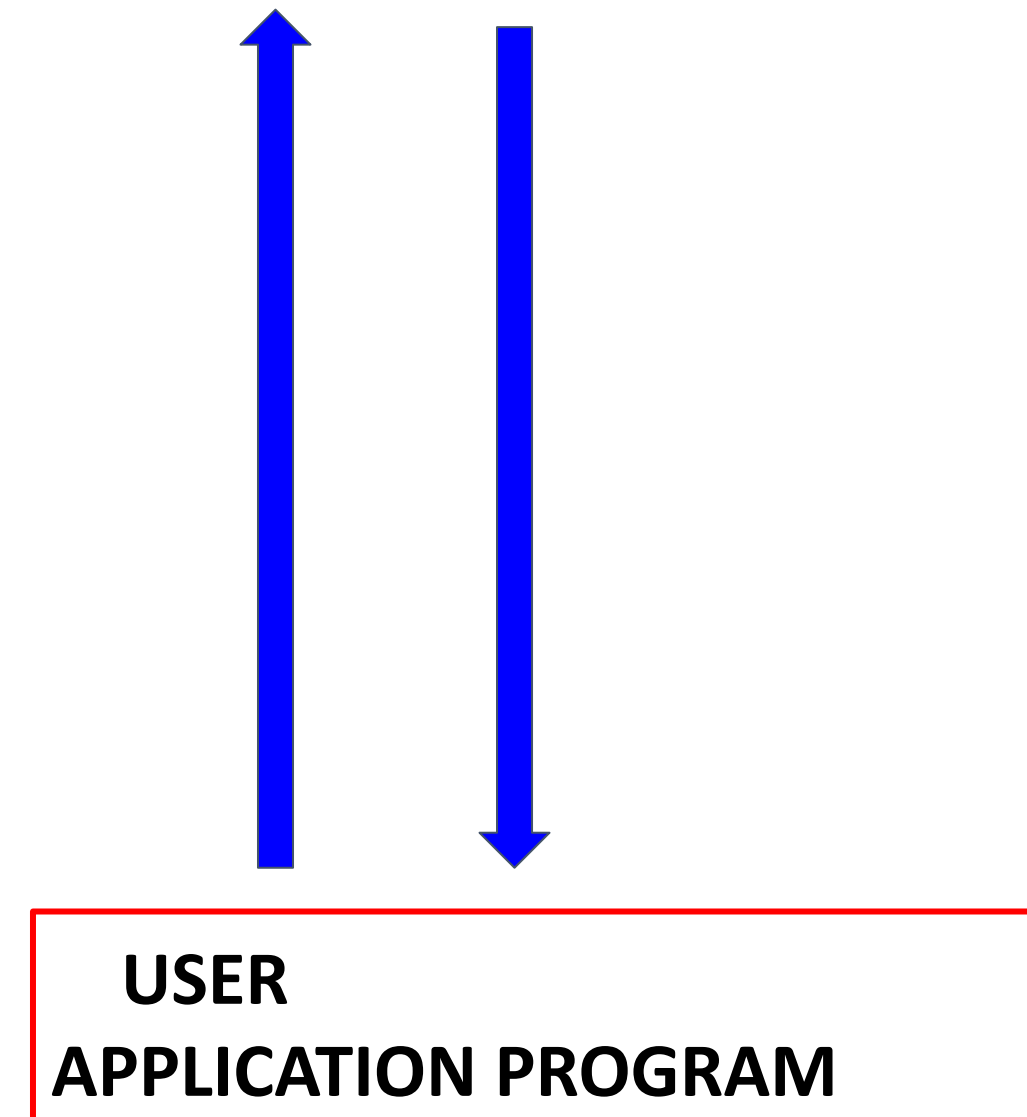


Using profiling tools like **gprof**

- Parallel programs performs well when atomic and synchronization constructs are reduced.
 - total removal not possible in all cases.
- Locality of reference importance for matrix operation
 - Spatial and Temporal Locality
- Access latency
 - In the increasing order: L1 cache, L2 cache, L3 cache, Main Memory, Hardisk (secondary storage)
- Sequential Consistency
 - Output of parallel program should be same as one sequential execution
 - Inter iteration Dependence (RAW/WAR/WAW) should not be there.
-

Operating System Structure

- Hardware at bottom
 - keyboard, mouse, network card etc.
- Operating system Kernel on top of it
 - Linux kernel
- Application Level Interface
 - Unix command line
- GUI interface
 - GNome



Today's Lecture: Profiling Parallel Programs

- Program
 - Need Space, Time, and Energy for computation
- Time - bash command
- gettimeofday(), and clock() functions.
- gprof
 - Performance Analysis Tool
 - Extended version of the tool prof
 - gprof can do callgraph collecting and printing


Parallel Matrix Addition

```
#pragma omp parallel for  
num_threads(8)
```

```
for(int i=0;i<size;i++)
```

```
for(int j=0;j<size;j++)
```

```
C[i][j]=A[i][j]+B[i][j];
```



Row Major Access

```
#pragma omp parallel for  
num_threads(8)
```

```
for(int i=0;i<size;i++)
```

```
for(int j=0;j<size;j++)
```

```
C[j][i]=A[j][i]+B[j][i];
```



Column Major Access

Sample Time Measurement Using time Command

```
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_row.c -fopenmp -o row
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./row 3

real    0m1.629s
user    0m2.067s ←
sys     0m0.337s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_column.c -fopenmp -o col
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./col 3

real    0m4.275s
user    0m21.490s ←
sys     0m0.565s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$
```

Time Command Output: Meanings

- Real
 - Time from start to finish of the execution.
- User
 - Amount of CPU time spend in user mode.
- Sys
 - Amount of time spend in kernel mode.

Sample time measurement using time command

```
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_row.c -fopenmp -o row
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./row 3

real    0m1.629s
user    0m2.067s
sys     0m0.337s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_column.c -fopenmp -o col
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./col 3

real    0m4.275s
user    0m21.490s
sys     0m0.565s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$
```

User time is total time taken by 8 threads.
The program was run using 8 threads.
Time per one core is around 1/8th.

Serial and Parallel Execution: Row Major addition

```
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_row.c -o row
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_row.c -o rowserial
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_row.c -fopenmp -o rowparallel
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./rowserial 3

TIME =373.191118
real    0m2.008s
user    0m1.747s
sys     0m0.260s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./rowparallel 3

TIME =126.043081
real    0m1.660s
user    0m2.178s
sys     0m0.352s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$
```

TIME value is in Milliseconds (Ms)
373 Ms = 0.373 Ms (**Serial**)
126 Ms = .126 Ms ($0.126 \times 8 = 1.008$) (**Parallel**)
TIME measured only for matrix addition loop shown in earlier slide.
Not for entire program.

Serial and Parallel Execution: Column Major addition

```
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_column.c -fopenmp -o colparallel
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ gcc matrixadd_column.c -o colserial
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./colparallel 3

TIME =2677.949905
real    0m4.271s
user    0m21.531s
sys     0m0.555s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$ time ./colserial 3

TIME =2193.084002
real    0m3.862s
user    0m3.511s
sys     0m0.348s
unnikrishnan@unnikrishnan-X510UNR:~/matrixadd$
```

TIME value is in Milliseconnds (Ms)
2668 Ms = (2668x8= 21344= 21.44 Seconds) (Parallel)
2193 Ms (Serial)
Time for Matrix addition loop shown in earlier slide.
Not for Entire program

times() function

- Signature : `clock_t times(struct tms *buf);`
- Returns the number of clock ticks elapsed.
- Stores the current process times in `buf`
- `struct tms {`
 - `clock_t tms_utime; // user time of process`
 - `clock_t tms_stime; //system time spend in kernel`
 - `....``}`

computing time using gettimeofday() function

- Function Signature
 - `int gettimeofday(struct timeval *, struct timezone *)`;
- `struct timeval`
 - stores number of seconds (`tv_sec`) and microseconds(`tv_usec`) since last Epoch (1970-01-01 00:00:00 +0000 -UTC)
- This can be used to get current relative time.
- To find time elapsed during execution of a code block
 - Get time immediately before entering the code block (`t1`)
 - Get time immediately after finishing execution of the code block (`t2`)
 - Value (`t2 - t1`) will give time elapsed during the execution.
 -

Value of “TIME” (see previous slide)

```
double rtclock()  
{  
    struct timezone Tzp;  
    struct timeval Tp;  
    int stat;  
    stat = gettimeofday (&Tp, &Tzp);  
    if (stat != 0) printf("Error");  
    return(Tp.tv_sec + Tp.tv_usec*1.0e-6);  
}
```

```
double t1= rtclock();  
    ..... /* Do Computation, matrix-add  
*/  
double t2= rtclock();  
double t3= (t2-t1);  
printf(“TIME=%f\n”, t3);
```


clock() function

- Signature - `clock_t clock(void);`
- Approximate processor time used by the program.

```
#include <time.h>
```

```
clock_t t1, t2;  
double cpu_time;
```

```
t1 = clock()  
... /* Do computation */  
t2 = clock();  
cpu_time= ((double) (t2 - t1)) / CLOCKS_PER_SEC; // Time in Seconds
```

CLOCKS_PER_SEC is having a value 1000000

GPROF - Usage

- Compile program with -pg option
 - gcc -pg program.c -o prog
- Run the program
 - ./prog [arguments]
 - generates gmon.out
- Run gprof with options
 - gprof [options] ./prog
 - it prints profiling information

Sample program

```
int fun1(){/*prg1.c begin */  
    printf("Inside fun1()\n");  
    for(int i=0;i<20000;i++);  
    return 1 ;  
/*prg1.c end */
```

```
int fun2() {/* prg.c begin */  
    printf("Inside fun2\n");  
    for(int i=0;i<1000000000;i++);  
    fun1();  
    return 1;  
}
```

```
static void fun3(){  
    printf("Inside fun3()\n");  
    for(int i=0;i<40000;i++);  
    return;  
}  
  
int main() {  
    printf("\n Inside main()");  
    for(int i=0;i<1000000000;i++);  
    fun2();  
    fun3();  
    return 0;
```

```
/*prg.c end */
```

Running the code for profile information

- Compile with profiling option
 - `gcc -pg prg.c prg1.c -o gprof_test`
- Run the program
 - `./gprof_test`
 - a file with name `gmon.out` will be created.
- Run with gprof tool
 - `gprof gprof_test gmon.out > output.txt`
 - `>` redirects output from screen to the file `output.txt`.

Flat profile

```
unnikrishnan@unnikrishnan-X510UNR: ~/gprof
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ ls
prg1.c  prg.c
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ gcc -pg prg.c prg1.c -o gprof_test
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ ./gprof_test
Inside main()
Inside fun2
Inside fun1()
Inside fun3()
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ ls
gmon.out  gprof_test  prg1.c  prg.c
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ gprof ./gprof_test gmon.out > output.txt
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ wc -l output.txt
157 output.txt
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ head -n 12 output.txt
Flat profile:

Each sample counts as 0.01 seconds.
 %   cumulative    self           self         total
time  seconds    seconds   calls  ms/call  ms/call  name
48.15      0.57      0.57           1    572.98    572.98  fun3
25.34      0.87      0.30           1    301.57    301.57  fun1
14.36      1.05      0.17           1     150.79    452.36  main
12.67      1.20      0.15           1     150.79    452.36  fun2

 %           the percentage of the total running time of the
time        program used by this function.
unnikrishnan@unnikrishnan-X510UNR:~/gprof$
```

command `head` with option `-n 12` prints first 12 lines of the file

total number of lines in output.txt is 157

Fields in the flat profile

- index - Unique value for each element in the table
- %time - Percentage of total time spent by the function
- Cumulative seconds
 - Total time of this function and those listed above in Seconds
- Self seconds
 - Number of seconds by this function alone
- calls
 - Number of times the function is called, if profiled. else blank
- Self/per call and total/per call -average value of all calls.
- name- Name of the current function

Flat profile

```
unnikrishnan@unnikrishnan-X510UNR: ~/gprof
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ ls
prg1.c  prg.c
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ gcc -pg prg.c prg1.c -o gprof_test
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ ./gprof_test
Inside main()
Inside fun2
Inside fun1()
Inside fun3()
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ ls
gmon.out  gprof_test  prg1.c  prg.c
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ gprof ./gprof_test gmon.out > output.txt
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ wc -l output.txt
157 output.txt
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ head -n 12 output.txt
Flat profile:

Each sample counts as 0.01 seconds.
 %   cumulative    self               self       total
time  seconds    seconds   calls  ms/call  ms/call  name
48.15      0.57      0.57           1    572.98    572.98  fun3
25.34      0.87      0.30           1    301.57    301.57  fun1
14.36      1.05      0.17           1     150.79    452.36  main
12.67      1.20      0.15           1     150.79    452.36  fun2

 %           the percentage of the total running time of the
time        program used by this function.
unnikrishnan@unnikrishnan-X510UNR:~/gprof$
```

command `head` with option `-n 12` prints first 12 lines of the file

total number of lines in output.txt is 157

Call graph profile

- index - unique id for each entry in the table
- %time
 - Percentage of total time spent by the function and its children
- Self seconds
 - Number of seconds by this function alone
- children
 - total amount of time propagated into this function by its children
- called - Number of times this function is called.
- name- name of the current function.
- <spontaneous>
 - if parent of function cannot be determined, this is printed for parent name.

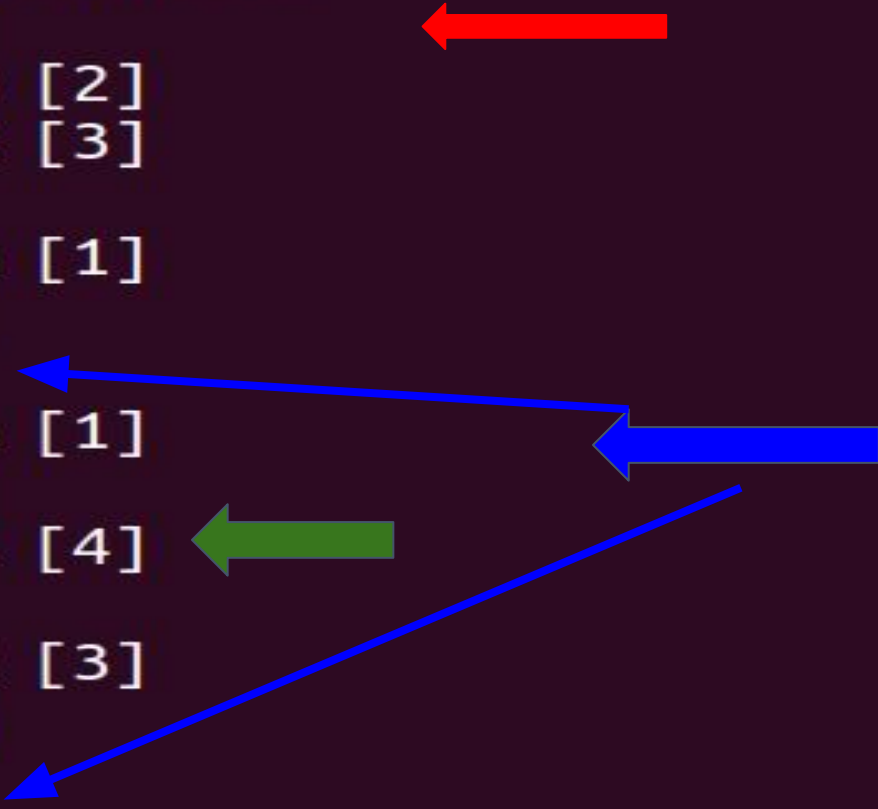
Call Graph - Profile information

```
unnikrishnan@unnikrishnan-X510UNR: ~/gprof
unnikrishnan@unnikrishnan-X510UNR:~/gprof$ head -n 65 output.txt | tail -n 22
Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 0.84% of 1.20 seconds

index % time      self  children  called  name
-----
[1]   100.0      0.17    1.03      1/1    main [1]
          0.57    0.00      1/1    fun3 [2]
          0.15    0.30      1/1    fun2 [3]
-----
[2]    47.9      0.57    0.00      1/1    main [1]
          0.57    0.00      1    fun3 [2]
-----
[3]    37.8      0.15    0.30      1/1    main [1]
          0.15    0.30      1    fun2 [3]
          0.30    0.00      1/1    fun1 [4]
-----
[4]    25.2      0.30    0.00      1/1    fun2 [3]
          0.30    0.00      1    fun1 [4]
-----

unnikrishnan@unnikrishnan-X510UNR:~/gprof$
```



The call graph diagram illustrates the relationships between the functions listed in the profile. A red arrow points from the 'main [1]' entry in the first row to the 'main [1]' entry in the second row. A blue arrow points from the 'main [1]' entry in the second row to the 'main [1]' entry in the third row. A green arrow points from the 'fun1 [4]' entry in the third row to the 'fun1 [4]' entry in the fourth row. Another blue arrow points from the 'fun1 [4]' entry in the third row to the 'fun1 [4]' entry in the fourth row.

gprof - reducing output text

- `gprof -b ./gprof_test gmon.out > output.txt`
 - removes the verbose information.
 - output was just 36 lines compared 157 without -b option.
- `gprof -p -b ./gprof_test gmon.out > output.txt`
 - prints only flat profile
 - output was just nine lines
- `gprof -pfun1 -b ./gprof_test gmon.out > output.txt`
 - prints flat profile only for fun1
- `gprof -P -b ./gprof_test gmon.out > output.txt` /* P - capital letter */
 - removes flat profile
- `gprof -q -b ./gprof_test gmon.out > output.txt`
 - prints only call graph profile.
- `gprof -qfun1 -b ./gprof_test gmon.out > output.txt` /*-Q to remove call graph profile */
 - prints call graph profile only for fun1

Sparse Matrix

Matrix with four rows and ten columns.
Non zero elements are twelve.

| | | | | | | | | | |
|---|----|----|----|----|---|---|----|----|----|
| 0 | 0 | 0 | 15 | 0 | 5 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 5 | 8 | 0 | 0 | 0 | 10 | 12 |
| 0 | 20 | 0 | 0 | 12 | 0 | 0 | 13 | 10 | 0 |
| 0 | 0 | 43 | 0 | 0 | 0 | 0 | 0 | 0 | 12 |

| | | | | |
|---|---|---|----|----|
| 0 | 2 | 6 | 10 | 12 |
|---|---|---|----|----|

| | | | | | | | | | | | |
|---|----|---|----|---|----|---|----|---|----|---|----|
| 3 | 15 | 5 | 5 | 3 | 5 | 4 | 8 | 8 | 10 | 9 | 12 |
| 1 | 20 | 4 | 12 | 7 | 13 | 8 | 10 | 2 | 43 | 9 | 12 |

Total Space in initial array - (4x10) elements
Modified array - (24+5)=29